

The BiolinX Affiliate Engine

A machine that finds creators already promoting competing peptide brands, proves they are real, decides who may be contacted, writes what to say at every turn of the conversation, and equips them the day they sign. Module by module, what runs on its own and what a person decides.

625

leads in the system

43

competitor brands
searched nightly

21,832

lines of code, 330
tests

\$0.14

cost per qualified lead
found

00 What you are looking at

The goal is 100 external affiliates by Black Friday. The method, in one line: find the creators who are already promoting a competitor for less than 25%, and offer them more. Everything in this document exists to make that repeatable instead of heroic.


HOW TO READ THIS DOCUMENT

Every module carries a badge: **AUTOMATIC** runs without anyone touching it, **HUMAN** waits for a person on purpose, **BY DESIGN** is something the system refuses to do, and **NEXT** is built but waiting on someone else. The refusals matter as much as the features; each one is explained where it appears.

Before this engine, recruiting ran through chat threads and a spreadsheet. Someone searched TikTok by hand, guessed whether a creator was real, wrote a message from scratch, sent it, and hoped they would remember to follow up. Nothing was counted, nothing was repeatable, and nothing survived the person doing it going on holiday.

What follows replaces that with thirteen modules that hand work to each other. Sourcing feeds verification, verification feeds the review desk, the review desk feeds outreach, outreach feeds sign-ups, sign-ups feed the asset generator. Every step writes down what it did and why, so any number in this document can be traced back to the post it came from.

The shape of the pipeline



Where it stands today

213 creators found by the engine	57 accepted, ready to contact	109 waiting for a human decision	312 leads tied to a named competitor	43 external affiliates in iDev
--	---	--	--	--

Plus 388 social handles matched to people, 509 enrichment records, 95 high-performing source posts banked for the content library, 625 job runs logged and 210 audited human actions. The engine has been running nightly without a miss since it was switched on.

Contents

01	Lead sourcing — the competitor-first discovery engine	how leads are found
02	Verification and the quality gates	how money is not wasted
03	Scoring, deduplication and ranking	who to work first
04	The review desk	the one human gate
05	The outreach flow	what to say, every turn
06	The outreach console	the seat an assistant works from
07	The compliance engine	the thing that keeps us legal
08	Sign-ups and provisioning	yes to account
09	The affiliate asset generator	what they post, written for them
10	The swipe file and content library	content that keeps coming
11	Affiliates, iDev and the goal counter	the number that matters
12	The admin, roles and the audit trail	who can do what
13	The scheduler and cost control	what runs, what it costs
14	What is automated and what is not	the master table
15	Why the first message is sent by a person	the platform rules
16	How it is built, tested and deployed	the engineering
17	What comes next, and what we need	the road ahead

01 Lead sourcing AUTOMATIC

The best affiliate is someone already promoting a competitor for less than we pay. So the engine does not search for peptide fans. It searches for competitors' discount codes, and collects the people using them.

How it works

Every night the engine takes the 43 active competitor brands and searches each one three ways on TikTok: "`<brand> code`", "`<brand> discount`", and the brand's domain. That is roughly 126 searches, run through a commercial scraping platform at about \$0.003 per result. Four audiences run in priority order: Weight-loss seeker, Biohacker, Gym / PED-curious, Anti-aging. Every audience is United States only, a rule that is never relaxed.

Hashtag searching was tried first and abandoned on evidence: a brand searched as a hashtag produced 10 leads, none of whom mentioned the brand. The same brand plus the word "code" returned real code posts immediately, 17 of 24. The engine keeps what worked.

The numbers

NIGHT	LEADS ADDED	SPENT	STOPPED BY
Wednesday 17th	16	\$3.15	audience spend caps
Thursday 18th	11	\$2.82	audience spend caps
Friday 19th	6	\$2.82	audience spend caps
Saturday 20th	11	\$2.82	audience spend caps

A larger one-off run on the 16th added 50 leads for \$6.90, touching 21 competitor brands in a night at \$0.14 per lead. Of those 50, 19 named a competitor *and* published a code, so they skipped review entirely and went straight into the outreach queue.

The things that make it cheap

- **Free checks come before paid ones.** A search result is only paid to be read if the post or bio already names a competitor. Last night that rule alone discarded 107 results for nothing.
- **Searches that stop working go to sleep.** Each search remembers how many new people it found. When a search stops producing, it rests and the budget moves to another brand. 40 of the searches are resting today, which is the engine telling us the competitor list needs widening.
- **Nothing is searched twice.** A search already run by another audience the same night is skipped.
- **Three spend caps:** per audience, per night, and a review-queue limit that stops sourcing when too many leads await a human.

What it captures per person

Handle, name, bio, followers, recent post statistics, the post that surfaced them, the competitor they named with their exact words, their code if they published one, an email if it is in their bio, and a country only when there is evidence for it.

02 Verification and quality gates AUTOMATIC

A follower count is not reach, and a profile is not a person. Everything the engine believes about a creator comes from a platform read, never from a guess.

The gates, in the order they run

GATE	WHAT IT CATCHES	WHEN
Names no competitor	A search result that mentions nobody we compete with	Free, before paying
Not English	Captions in another language	Free
Off-niche	Hair oil, art, unrelated content	Free
Follower range	Too small to matter, too big to answer	Free, rechecked after the read
Outside the US	Platform field, post country, bio wording, agreeing captions	After the read
Dead account	No post in 60 days	After the read
Weak reach	Median views under 0.5% of followers on TikTok, 1% on YouTube	After the read
A shop, not a person	Storefront or clinic wording, business category	After the read
The brand's own page	@brandnameofficial, @brandlabs001, @brandfulfillment	After the read

What the gates dropped last night

107 named no competitor	2 outside the US	3 brand's own page	1 dead account	11 saved as leads
-----------------------------------	----------------------------	------------------------------	--------------------------	-----------------------------

WEAK REACH: THE GATE THAT PAID FOR ITSELF

An independent review of 47 leads found the main failure was creators with large follower counts and almost no views — one had 466,000 followers and 526 views a post. Bought or dead audiences. The engine now measures median views against followers and rejects the mismatch. When that rule was applied backwards over the existing queue it rejected 15 leads that a person had already approved.

It also taught us to be careful: the first version of the rule rejected a real creator because the platform reported zero views on some posts. Zero-view rows are now ignored as unreported, and that creator was restored with a note in the audit log explaining why.

Nothing is paid for twice

A creator who fails a gate is remembered for 90 days, so the engine never pays to read them twice. A failed read is not remembered, because a timeout is not a verdict. Every saved creator keeps the record of their read, so "why did we think she was in the US?" always has an answer.

03 Scoring, deduplication and ranking AUTOMATIC

Two questions: is this the same person we already have, and who should be contacted first? Both are answered by rules, not by a model, so both can be explained.

Deduplication on five keys

The same creator can arrive as a TikTok handle today, an email in a bio next week, and a name on a spreadsheet next month. Five keys are checked against every table that holds a person: discount code, platform handle, email, website, and name plus platform. Handles are parsed out of every format the imported data used. A unique database index on handles is the final backstop, inside a transaction, so two runs at the same moment cannot both create the same person.

READING A CODE CORRECTLY IS HARDER THAN IT SOUNDS

Early on the engine saved ordinary words as discount codes — FOR, THEM, NONE, SHARE — and then treated those creators as already known, silently blocking real affiliates. Now a code is only read from "use code X", a quoted code, a disguised spelling like "cod3", or a referral link on that competitor's own domain. Stop words are rejected, and when a post names two brands the code nearest the mention wins.

Scoring, with a reason for every point

SIGNAL	WHY IT MATTERS
Already an affiliate of a named competitor	The strongest signal there is. They promote peptides for money today
Their competitor's rate against our 25%	Decides which offer they get
Verified reach inside the range	Real audience, reachable size
Active in the last 30 days	Dormant accounts do not post for us either
Track record of promotions	They have done this before, and their audience accepts it
Original content, on-niche words	Repost accounts convert badly
Store or business account	A penalty: a shop is not a creator

Scores become bands, bands become a conversion rank, and the rank sets the order of the work queue. A lead waiting for review has no rank at all, which is how the system guarantees an unreviewed creator is never contacted by accident.

WHY THE SCORING IS NOT AN AI JUDGEMENT

Every point has a written reason attached, visible in the admin. When marketing disagrees with a score they can see the sentence that caused it and change the rule. A model that says "8 out of 10" cannot be argued with, cannot be tested, and cannot be audited two months later. Claude is used in this system for language work — writing, summarising, classifying replies — and never for the decisions that money depends on.

04 The review desk HUMAN

One human gate, placed where judgement actually helps: deciding whether a creator is worth the team's time, with all the evidence on one screen.

What the reviewer sees

Each sourced lead arrives with its handle, bio, niche, country with a flag for unknown, reach, average views, engagement rate, posts in the last 30 days, the competitor named with the exact quoted words, their code, a store flag, the surfaced post and its statistics, and the score with its reasons. Three clicks: accept, reject with a reason, or open the full profile. Bulk accept and reject handle a run of obvious ones, with a confirmation step.

The one case that skips review

A creator who names a competitor *and* publishes a personal discount code is, by definition, that competitor's affiliate. There is nothing for a human to judge. Those are accepted automatically and go straight into the outreach queue. A brand's own account is dropped instead, and anything without a code still waits for a person. That rule turned 19 of the last 50 finds into immediately workable leads.

The protected category

BY DESIGN GOODWILL ADVOCATES ARE NEVER CONTACTED

Some creators already speak warmly about BiolinX without any commercial relationship. A recruitment pitch to them is not a missed opportunity, it is damage. They are marked as a protected category, and the system refuses to draft, approve or record a message to them anywhere in the code. A reviewer cannot classify someone into that category by accident either; it can only be set deliberately.

Why the queue has a ceiling

Sourcing stops adding leads once too many are waiting for review — the limit is set to 150 today. The bottleneck in this business is human attention, not scraping capacity, and a queue nobody can clear is money spent for nothing. When the queue empties, sourcing resumes on its own.

57

accepted

109

waiting

47

rejected, with reasons

150

queue ceiling

05 The outreach flow AUTOMATIC

The conversation is a decision tree, not a template. The system knows which message is next for every creator, at every turn, and nobody on the team has to remember any of it.

Which offer a creator gets

SITUATION	WHAT HAPPENS
Their competitor pays under 25%, or the rate is unknown	Offer 1 — we beat it, 25% for life
Their competitor already pays 25%	Offer 2 — match it and win on the perks
Their competitor pays more than 25%	No message. The flow has no offer that wins
They promote a brand we cannot identify	No message until a person names the brand
Not signed with anyone	No message. This programme recruits competitor affiliates

Every branch of the conversation

Each opener comes in a soft and a direct version, split evenly across creators so the two can be compared, and opens with one of seven rotating curiosity lines so two creators in the same niche never receive an identical message. From there the tree covers every answer:

- **Yes** → ask for the four sign-up details
- **Tell me more** → send the programme details, then the sign-up ask
- **No** to Offer 1 → the recruitment commission and the sister brand, Aro
- **No** again → a warm referral ask, and the conversation closes politely
- **Replied but said nothing useful** → the details, not a sign-up ask they never agreed to
- **No reply** → a check-in after the wait, a maximum of two, then closed

Thirteen message templates in total. Marketing owns the words and edits them in the admin without a developer; the flow owns when each one is used. The wording written by the marketing team is marked as approved copy, which exempts it from two rules meant for machine-written cold messages while every other compliance rule still applies to it.

WHAT THE SYSTEM WILL NOT LET HAPPEN

- The same message recorded twice, even if the button is pressed twice or an old tab is left open
- A message out of order — only the message the flow expects now can be recorded
- A reply logged before anything was sent
- A second sign-up for the same creator
- A creator contacted from two places at once; the older queue and the new flow know about each other
- A message to anyone unqualified, protected, dead, or already signed

06 The outreach console HUMAN

A seat built so the person sending messages needs no training, no judgement and no knowledge of the system. One lead at a time, the message already written, one button to move on.

What the assistant sees

The queue decides the order: replies to answer first, then check-ins due, then new creators, with the ones whose profile we can open directly before the ones needing a name search. The card shows who they are, how many followers, which competitor they promote and the exact words they used, then the message, already filled in.

1. **Warm them up** — a three-tick checklist: follow, like two posts, leave one real comment. Twenty seconds, and a message from an account the creator just saw in their notifications gets opened instead of ignored.
2. **Copy the message** — one button, with a fallback if the browser blocks it.
3. **Open the chat** — for Instagram this opens the message thread directly rather than the profile. Their profile stays one click away so the assistant can check it is the right person.
4. **I sent it** — records the message, sets the follow-up date, loads the next creator.
5. **Paste the reply** — the answer is classified, and the next message appears underneath.
6. **Save the sign-up** — the form arrives pre-filled from what the creator wrote.

Reading replies

A pasted reply is classified by keyword rules first — yes, tell me more, no, or replied without details — and the matching button is pre-selected. When the keywords are not confident, and only then, a small language model reads it and suggests. The assistant can always overrule with one click. Names, emails and requested codes are pulled out of the message text and carried into the sign-up form.

DETAILS THAT ONLY MATTER WHEN A REAL PERSON SITS THERE ALL DAY

- A creator handed to one assistant is held for fifteen minutes so two people never message the same person
- Skipping asks why: the account is gone, or they are not a fit, and the answer is recorded
- Edited text is re-checked against compliance before the send button unlocks
- The whole page works on a phone, with the action button fixed to the bottom of the screen
- Assistants get an account that can only see this page and sign-ups, never the lead database

07 The compliance engine BY DESIGN

BiolinX sells research-use-only products. One careless sentence in one message is a regulatory problem, not a marketing mistake. Every outbound string in this system passes the same linter before a human is allowed to send it.

The rules

RULE	WHAT IT BLOCKS
Banned compounds	Brand drug names and regulated compound names, including disguised spellings
Claims	Weight, muscle, aging, dosing, protocols, treatment, cures, prescriptions, "clinically proven"
Personal use	"I take it", "works for me", "my results" — nobody may imply human use of a research product
Message shape	Dashes, over-long openers, a commission figure in a first message, links in a first message
Restricted products	The restricted catalogue, named or skeleton-matched
Research-use line	Any public content must carry it, word for word
Email law	Unsubscribe link, postal address, US recipients, never a scraped address, never a suppressed one
Earnings	Any figure about money earned must sit next to a historical qualifier

Written to survive people trying to get around it

Text is checked three ways at once: as written, with lookalike characters folded back to plain letters, and as a letters-only skeleton that defeats spacing, leetspeak and homoglyph tricks. That work was not theoretical — a red-team pass found each of those evasions before the rules were hardened.

WHAT THIS COSTS US, HONESTLY

Competitors' affiliate posts are openly non-compliant, and they stay up. Our affiliate assets read more plainly than theirs because they must. The marketing team's own draft templates included "where I get mine" and "my source", and both are blocked here — they imply personal use of a research product. That is a deliberate trade: slightly less punchy copy, no regulatory exposure, and no affiliate of ours is ever handed wording that could harm them.

08 Sign-ups and provisioning

AUTOMATIC

HUMAN

From "yes" to a working account. Recorded in one place, checked, and carried through the steps that create a real affiliate.

What happens on yes

The assistant saves four details — first name, last name, email, and the code the creator wants. The system links the sign-up to the lead it came from, credits the recruiter as a real database reference rather than a typed name, marks the lead as signed, closes the outreach conversation, and puts the sign-up on the Sign-ups page.

Provisioning is built as a saga: each step is stored separately, so a failure halfway through is visible and repairable instead of leaving a half-made affiliate. The steps are the coupon in the store, the coupon assigned in iDev, the welcome email, the tax form, the payout authorisation, and the final confirmation by the person who owns the affiliate admin.

NEXT

THE ONE PIECE THAT IS STILL MANUAL

iDev's API hands back an affiliate's name, email and sign-up date, but not their discount code, and there is no way to create a code from outside the store admin. So today a person checks the code is free, creates it, and assigns it, in two systems.

A specification has been written for four small endpoints on the BiolinX side: list codes, check a code is free, create one, and report what each code has sold. With the first, every affiliate's assets fill themselves in. With the next two, the whole path from "yes" to a working code runs without a person. The fourth finally answers which affiliates actually sell.

Customer.io

Every lead with an email is mirrored as a person in Customer.io hourly and on acceptance, carrying the attributes the marketing specification asked for. Suppressed addresses are sent as unsubscribed and nothing else. No campaign is ever triggered from here; the engine supplies the audience, marketing owns the sending.

09 The affiliate asset generator AUTOMATIC

The day someone signs, they get ten pieces of copy with their own code already in them. No affiliate should have to invent the wording, and no affiliate should invent it badly.

What is generated

ASSET	WHERE IT GOES
Bio line	Instagram or TikTok bio, fitted to the 150-character limit
Link title	Linktree or the link in their bio
Four story angles	Curiosity, code, partner, partner plus code
Feed caption	Under a TikTok or Reel, with the paid-partnership disclosure
Ad-break script	Twenty seconds to read out in a YouTube video
On-screen banner	The strip along the bottom while they read it
Description block	Top of the video description, with the affiliate disclosure

Each one shows where it goes, counts its characters against what that surface allows, and is checked by the same compliance linter that guards our own messages — as public content, which means the research-use line is mandatory on every single one. A copy button per asset, and the button stays locked if anything is missing or blocked.

THE BIO LINE, AS AN EXAMPLE OF THE CARE INVOLVED

Instagram allows 150 characters in a bio. The research-use line takes 94 of them, and it cannot be shortened without breaking the rule it exists for. That leaves 56 characters to name the brand, the code and the discount. The generated line comes out at 140 characters, checked automatically, so no affiliate ever pastes something that gets cut off halfway through the legal sentence.

10 The swipe file and content library AUTOMATIC HUMAN

Affiliates promote more when there is something ready to post. This module studies what already works in the niche and turns it into original BiolinX content, without copying anyone.

How a post is made

1. **Find what worked.** Posts that beat their own creator's average views by half again, with at least ten thousand views, recent, in English. Plus a separate search of the niche's top posts, which banked 95 source posts for about a dollar.
2. **Write something new.** Claude writes three original variants from the idea, never the words: a hook, a caption placing the discount code once, hashtags, the words for the image and a brief describing what the image should show.
3. **Check before anyone sees it.** Every variant is run against BiolinX's own content rules and our compliance linter. The best passing variant becomes a draft; if none pass, the writer repairs it once with the exact reasons.
4. **A person decides.** BiolinX checks compliance, not quality, so quality is judged here. Approve, or decline with a note saying what is missing — and a new version is written from that note.
5. **Hand over.** Approved posts are sent signed to the BiolinX content library, where they wait as drafts until someone presses Publish. Their image generator makes the picture from our brief and calls back with the link.


Both directions are signed and time-limited, callbacks are verified against the exact bytes received, and the whole exchange is documented so the two teams can work without meetings. 95 source posts are banked and waiting today.

11 Affiliates, iDev and the goal counter AUTOMATIC

The number the whole project is judged on is 100 external affiliates. That number is computed, not typed, and it is deliberately conservative.

The affiliate platform is read every thirty minutes. Its API is fragile in specific ways — it returns an empty response above a certain page size, and certain records break its own encoder — so the reader handles both: a fixed page size, and a bisecting retry that isolates a broken record, records exactly which position failed, and carries on. A broken record is surfaced, never silently dropped.

Every affiliate is then classified as a real partner, one of us, or undecided. The engine suggests which are team and house accounts; a person confirms. The dashboard counts confirmed partners only, and shows a range while anything is still undecided, so the headline number is never flattering by accident.



The referral tier that pays affiliates for recruiting other affiliates is read and monitored, and the setting that controls it is never changed by the engine.

12 The admin, roles and the audit trail HUMAN

One place for the whole team, with each person seeing exactly what their job needs and nothing else.

The pages

PAGE	WHAT IT IS FOR
Dashboard	The goal, the funnel end to end, the last job runs
Outreach	The assistant's seat: one creator at a time
Leads	Queue, sourced, all, needs triage — with filters, sorting and CSV export
Audiences	Where to search, how much to spend, the competitor list
Message templates	Marketing edits every word of the outreach flow, with live previews
Swipe file	Approve or decline content before it reaches affiliates
Send messages, Email ops, Replies, Sign-ups, Affiliates	The rest of the pipeline
Activity	Every job run and every human action, in order
Team, Settings	Invites, roles, and every key stored encrypted

Roles

Four roles: administrator, operations, representative and outreach assistant. The assistant role sees only the outreach page and sign-ups, and is blocked from the lead database at the server, not just hidden in the interface. Invitations are links; no passwords are ever emailed.

Everything is recorded

210 human actions are in the audit log so far: who accepted which lead, who rejected one and why, who approved a message, who changed a setting. When a correction is applied in bulk — as when the weak-reach rule was applied backwards over existing leads — it is recorded as a job with its reason, so the history explains itself months later.

13 The scheduler and cost control AUTOMATIC

A server that spends money needs rules about when it is allowed to. These are those rules.

JOB	HOW OFTEN	SPENDS	WHAT IT DOES
Affiliate sync	30 minutes	No	Reads the affiliate platform, updates the goal count
Rank recompute	6 hours	No	Re-scores and reorders the queue
Referral expiry	Daily	No	Watches referral windows
Customer.io sync	Hourly	No	Mirrors leads with emails
Swipe sync	10 minutes	No	Checks the content library for status changes
Lead sourcing	Daily, 23:37	Yes	The nightly search of every competitor brand
Research and personalise	Daily	Yes	Reads accepted leads in depth, writes talking points
Content search	Daily when switched on	Yes	Tops up the swipe file when it runs low

THE FOUR RULES THAT KEEP SPENDING SANE

1. **Never on boot.** A restart must not spend money. Jobs that cost anything wait at least an hour and are anchored to their last real run.
2. **Three caps.** Per audience, per day across everything, and a hard stop when the review queue is full.
3. **Only one server runs a job.** A database lock means two processes never do the same paid work twice.
4. **Every run is recorded** with what it found and what it cost — 625 runs so far, each one inspectable.

Total spend on discovery since the engine was switched on is under thirty dollars, for 213 verified creators, their competitor evidence and their contact details.

14 What is automated, and what is not

The honest master table. Where a person is needed, the reason is in the last column — and in every case it is a deliberate decision, not an unfinished feature.

THE SHORT VERSION

Everything that can be automated without breaking a platform rule or a regulation, is. What is left for a person is three things: deciding whether a creator is worth pursuing, pressing send, and judging whether content is good. Those are the three places where human judgement is the product, not the bottleneck.

STEP	WHO DOES IT	WHY
Finding creators promoting competitors	AUTOMATIC	Nightly, 43 brands, inside a spend cap
Proving reach, country and activity	AUTOMATIC	Read from the platform, never estimated
Reading their competitor and code	AUTOMATIC	From their own words, with the quote kept
Scoring, deduplicating, ranking	AUTOMATIC	Rules with reasons, not a model
Accepting or rejecting a lead	HUMAN	Judgement about people is worth a person's minute
Accepting a creator who published a code	AUTOMATIC	The evidence is conclusive; there is nothing to judge
Choosing which message comes next	AUTOMATIC	The flow chart, applied the same way every time
Writing the message	AUTOMATIC	Marketing's words, filled in per creator
Checking it is legal to send	AUTOMATIC	Every character, every time
Sending the first message	HUMAN	No platform permits software to do this. See section 15
Classifying the reply	AUTOMATIC	Keywords first, a model only when unsure
Deciding the answer to a reply	AUTOMATIC	The flow chart again
Sending that answer	HUMAN	Same platform rule
Scheduling follow-ups and check-ins	AUTOMATIC	Dates set on send, queue ordered by urgency
Recording the sign-up	AUTOMATIC	Pre-filled from the creator's own message
Creating their discount code	HUMAN NEXT	No API exists yet; specification written and sent
Generating what they post	AUTOMATIC	Ten assets, compliance-checked
Writing content for the library	AUTOMATIC	Three original variants per idea
Approving that content	HUMAN	Compliance is machine-checked; quality is not

STEP	WHO DOES IT	WHY
Counting affiliates against the goal	AUTOMATIC	Read from the affiliate platform every 30 minutes
Deciding a person is team, not partner	HUMAN	Suggested by the engine, confirmed by a person
Checking affiliates actually post	NEXT	Waiting on the posting rules to be agreed

15 Why the first message is sent by a person

This is the question every client asks, so here is the full answer with the evidence, because the wrong decision here costs the brand account rather than a few hours.

The platforms, one by one

PLATFORM	CAN SOFTWARE SEND A FIRST MESSAGE?	WHAT IS ALLOWED
Instagram	No. The messaging API only replies to people who message first, within 24 hours	Automatic replies, and a message to someone who comments on our post
TikTok	No. Its business messaging API states conversations cannot be initiated	Replies within 48 hours; the comment trigger is not available in the United States
YouTube	Direct messages do not exist	Email, where it is published
Reddit	Technically possible, and unsolicited bulk messages get the account suspended	Nothing safe at volume

What the automation tools actually are

Tools that promise automatic direct messages do not use any official interface. They drive the app the way a person would, from many accounts at once. One such tool advertises, on its own sales page, "IP change methods to avoid account lockouts". A product that sells a way to dodge lockouts is telling you plainly that it causes them.

THE ARGUMENT THAT SETTLES IT

A bot does not send more messages than a person. Both platforms cap how many messages any account may send in a day, and the cap does not care what is typing. So automation buys no extra volume — it buys the same volume with the loss of the brand account attached. For a research peptide brand, that account is not replaceable.

There is a second cost people miss. A message from an account the recipient does not follow lands in a requests folder that most creators never open. The reply rate collapses before any block occurs. And the people being messaged are competitor affiliates in one small niche who talk to each other; a spam blast burns the exact list the engine spent money building.

So the system automates everything up to the send

Finding, verifying, qualifying, choosing the message, writing it, checking it, scheduling the follow-up, reading the answer and deciding the reply — all machine. The person taps copy, pastes, and presses one button, about a minute per creator. One assistant covers roughly forty creators a day; two cover four hundred a week, with no risk to the account.

Where automation *is* permitted, it is on the list to build: automatic replies to people who message us first, and the comment-to-message flow on Instagram, which is allowed in the United States and needs only a business account and an approval from the platform.

16 How it is built, tested and deployed

Nineteen days from first commit to a system that runs itself nightly. What is under the bonnet, and why it can be trusted with money.

21,832

lines of code

153

source files, 8
modules

330

automated tests

90

commits since 2
September

625

job runs, all recorded

The pieces

MODULE	WHAT LIVES THERE
core	Ranking, cadence, the outreach flow chart, message templates, asset templates, settings
compliance	The linter, the banned lists, the evasion-resistant matching
scraping	Platform readers, discovery, scoring, quality gates, country evidence, deduplication
jobs	Everything scheduled: sourcing, enrichment, outreach, sign-ups, content, sync
db	The schema and its migrations, settings encrypted at rest
api / admin / worker	The server, the interface, and the process that runs the schedule

How quality is kept

- **Tests come with the rules.** 330 of them, each one written from a real case — the ordinary word saved as a discount code, the brand page hiding behind a number, the creator wrongly rejected for zero-view posts.
- **Three independent review passes** have been run over the system by separate reviewers, plus two end-to-end tests against the live database that were rolled back afterwards. Roughly thirty defects were found and fixed before anyone met them.
- **Every deploy is one command** that pulls, installs, applies database changes, builds and restarts, then waits for the health check before reporting success.
- **Secrets never touch the code.** Keys live encrypted in the database or in the server environment, and are redacted from every log line.

THINGS LEARNED THE HARD WAY, NOW WRITTEN DOWN

Scraped text needs whole-word matching, because squashing it made "information peptides" look like a brand called "Ion Peptide". The affiliate platform returns an empty body rather than an error for certain records. A creator's dates arrive in six different formats across platforms. A long job started from a terminal session dies when the session ends. Every one of these is documented in the repository so the next person does not rediscover it.

17 What comes next, and what we need

The engine is running. These are the things that raise its output, in the order they matter.

Ready to build, waiting on a decision

WHAT	WAITING ON	WHAT IT UNLOCKS
Posting rules and monitoring	The rules themselves: how many posts a month, and what happens if they do not	Turns "use it or lose it" into a report, using the readers we already have
Partner code endpoints	The BiolinX development side; specification already written	Every affiliate's assets fill themselves in; sign-up becomes fully automatic
Aro programme details	Commission, cookie length and a sign-up link	Three messages in the outreach flow cannot be sent without them
Automatic replies and comment-to-message	A business account on the social platforms	The automation that platforms do permit
Email to published addresses	A decision on whether we email addresses creators publish	The one outreach channel that can run end to end without a person
Alerting	A messaging token	A failed job tells someone instead of waiting to be noticed

The two things that most raise output right now

MORE COMPETITOR BRANDS

Each brand added is three new searches. Of the searches we have, 40 are now resting because they have found everyone they are going to find. The competitor list is the engine's fuel, and it is the cheapest thing on this page to top up.

SOMEBODY TO CLEAR THE REVIEW QUEUE

109 creators are waiting for a yes or no, and sourcing stops itself at 150 on purpose. Each decision takes seconds with the evidence on screen. This is the difference between a system that is running and a system that is producing.

The scale it is built for

The design target was five thousand scored, deduplicated leads and a hundred external affiliates. The limits in place today are deliberately conservative: a ten dollar daily spend, a hundred and fifty lead review queue, four audiences. Every one of them is a number in a settings page, not a rewrite. The work of the next phase is not more engineering; it is more brands, more reviewers, and more assistants in a seat that already exists.

IN ONE PARAGRAPH

The engine finds creators who already sell peptides for someone else, proves they are real, decides who is worth contacting, writes the right message for every branch of the conversation, checks every word against the rules that keep a research-use brand safe, hands it to a person to send in about a minute, reads the answer, schedules the follow-up, records the sign-up, and hands the new affiliate ten pieces of copy with their own code in them. It costs about fourteen cents to find a qualified creator, and it has never missed a night.

BiolinX Affiliate Engine · system documentation · 21 September 2026